



Training dynamics — 4x60 tanh, hard IC, 40474 uniform points, 27 windows, causal


(a) Total loss convergence



(b) Loss decomposition (hard IC: no penalty term)



(c) Learning-rate schedule



(d) Logged loss vs reference error (coverage may grow)

